

ASSIGNMENT-3

Assignment Date	8 October 2022
Student Name	Ms.Brundha B
Student Roll Number	192IT128
Maximum Marks	2 marks
Team ID	PNT2022TMID01939

```
In [1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns

from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
import statsmodels.api as sma
from statsmodels.stats.outliers_influence import variance_inflation_factor

from sklearn.linear_model import LogisticRegression
from sklearn.svm import SVC
from sklearn.neighbors import KNeighborsClassifier
from sklearn.tree import DecisionTreeClassifier
from sklearn.model_selection import cross_val_score

from sklearn.model_selection import GridSearchCV
from sklearn.metrics import classification_report
from sklearn.metrics import confusion_matrix
import warnings
```

```
warnings.filterwarnings('ignore')
```


```
In [2]: df = pd.read_csv("/content/drive/MyDrive/abalone.csv")
```

```
In [3]: df.rename(columns={"Sex": "sex", "Length": "length", "Diameter": "diameter",
                        "Height": "height", "Whole weight": "whole_weight",
                        "Shucked weight": "shucked_weight", "Viscera weight": "viscera",
                        "Shell weight": "shell_weight", "Rings": "rings"}, inplace = True)
```

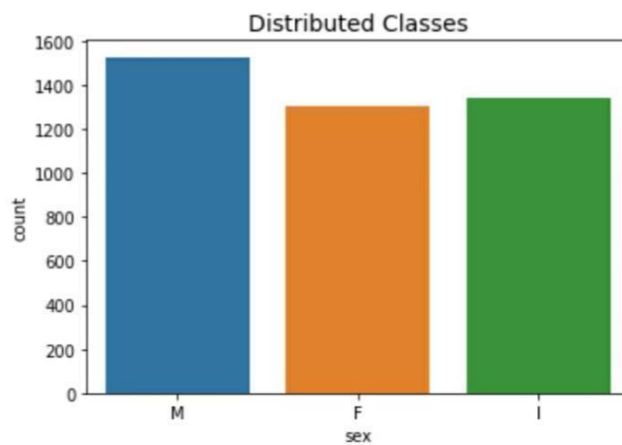
```
In [4]: df[df['height'] == 0] #need to drop these rows.
df.drop(index=[1257,3996], inplace = True)
df.shape
```

Out[4]: (4175, 9)

```
In [5]: df['age'] = df['rings'] + 1.5 #AS per the problem statement
df.head()
```

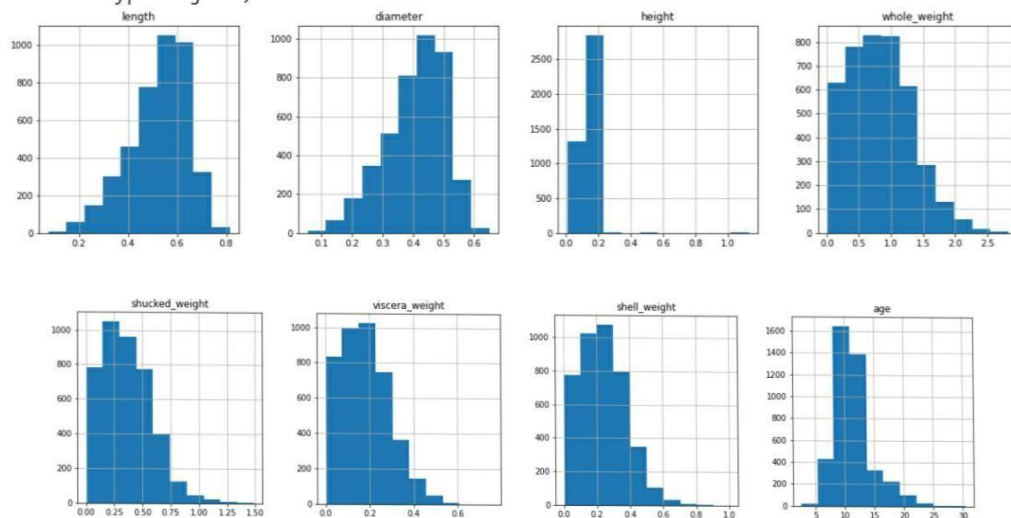


```
In [6]: sns.countplot('sex', data=df)
plt.title('Distributed Classes', fontsize=14)
plt.show()
```



```
In [7]: df.hist(figsize = (20,10), layout = (2,4))
```


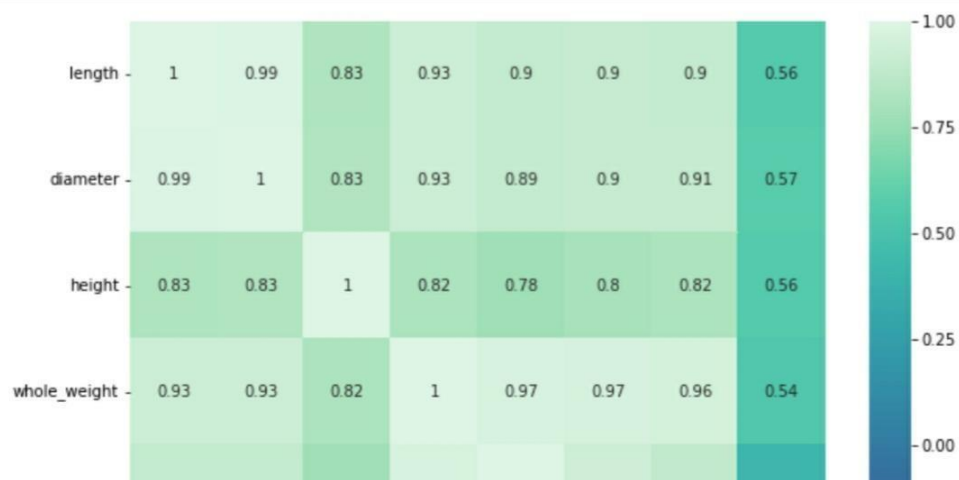
```
Out[7]: array([[<matplotlib.axes._subplots.AxesSubplot object at 0x7f9818f073d0>,
<matplotlib.axes._subplots.AxesSubplot object at 0x7f9818ebc990>,
<matplotlib.axes._subplots.AxesSubplot object at 0x7f9818ef3f90>,
<matplotlib.axes._subplots.AxesSubplot object at 0x7f9818eb55d0>],
[<matplotlib.axes._subplots.AxesSubplot object at 0x7f9818e6dbd0>,
<matplotlib.axes._subplots.AxesSubplot object at 0x7f9818ecadd0>,
<matplotlib.axes._subplots.AxesSubplot object at 0x7f9818f850d0>,
<matplotlib.axes._subplots.AxesSubplot object at 0x7f9819481250>]],
dtype=object)
```



```
In [8]: df.skew().sort_values(ascending = False)
```

```
Out[8]: height      3.166364
age      1.113754
shucked_weight  0.718735
shell_weight   0.621081
viscera_weight  0.591455
whole_weight   0.530549
diameter     -0.610182
length       -0.640993
dtype: float64
```

```
In [9]: corr = df.corr()
plt.figure(figsize = (10,10))
ax = sns.heatmap(corr, vmin = -1, center = 0, annot = True, cmap = 'mako')
```



```
In [10]: upper_tri = corr.where(np.triu(np.ones(corr.shape),k=1).astype(np.bool))
columns_to_drop = [column for column in upper_tri.columns if any(upper_tri[column, column] == 0)]
print("Columns to drop:\n", columns_to_drop)
```

```
Columns to drop:
['diameter', 'shucked_weight', 'viscera_weight', 'shell_weight']
```

```
In [11]: df.head()
```

```
Out[11]:
```

	sex	length	diameter	height	whole_weight	shucked_weight	viscera_weight	shell_weight	age
0	M	0.455	0.365	0.095	0.5140	0.2245	0.1010	0.150	16.5
1	M	0.350	0.265	0.090	0.2255	0.0995	0.0485	0.070	8.5
2	F	0.530	0.420	0.135	0.6770	0.2565	0.1415	0.210	10.5
3	M	0.440	0.365	0.125	0.5160	0.2155	0.1140	0.155	11.5
4	I	0.330	0.255	0.080	0.2050	0.0895	0.0395	0.055	8.5

```
In [12]: df.shape
```

```
Out[12]: (4175, 9)
```

```
In [13]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 4175 entries, 0 to 4176
Data columns (total 9 columns):
#   Column          Non-Null Count  Dtype
---  ---
0   sex              4175 non-null   object
1   length           4175 non-null   float64
2   diameter         4175 non-null   float64
3   height           4175 non-null   float64
4   whole_weight     4175 non-null   float64
5   shucked_weight   4175 non-null   float64
6   viscera_weight   4175 non-null   float64
7   shell_weight     4175 non-null   float64
8   age              4175 non-null   float64
dtypes: float64(8), object(1)
```

```
memory usage: 455.2+ KB
```

```
In [14]: df[df.duplicated()]
```


```
Out[14]: sex length diameter height whole_weight shucked_weight viscera_weight shell_weight age
```





```
In [15]: for i in df:
          if df[i].dtype=='int64' or df[i].dtypes=='float64':
              q1=df[i].quantile(0.25)
              q3=df[i].quantile(0.75)
              iqr=q3-q1
              upper=q3+1.5*iqr
              lower=q1-1.5*iqr
              df[i]=np.where(df[i] >upper, upper, df[i])
              df[i]=np.where(df[i] <lower, lower, df[i])
```


```
In [16]: import matplotlib.pyplot as mtp
```

```
In [17]: def box_scatter(data, x, y):
          fig, (ax1, ax2) = plt.subplots(nrows=2, ncols=1, figsize=(16,6))
          sns.boxplot(data=data, x=x, ax=ax1)
          sns.scatterplot(data=data, x=x,y=y,ax=ax2)
```

```
In [18]: for i in df:
          if df[i].dtype=='int64' or df[i].dtypes=='float64':
              mtp.boxplot(df[i])
```







In [19]: `df.head()`

Out[19]:

	sex	length	diameter	height	whole_weight	shucked_weight	viscera_weight	shell_weight	age
0	M	0.455	0.365	0.095	0.5140	0.2245	0.1010	0.150	16.5
1	M	0.350	0.265	0.090	0.2255	0.0995	0.0485	0.070	8.5
2	F	0.530	0.420	0.135	0.6770	0.2565	0.1415	0.210	10.5
3	M	0.440	0.365	0.125	0.5160	0.2155	0.1140	0.155	11.5
4	I	0.330	0.255	0.080	0.2050	0.0895	0.0395	0.055	8.5

In [20]:

```

from sklearn.preprocessing import LabelEncoder
encoder=LabelEncoder()
df['sex']=encoder.fit_transform(df['sex'])
df.head()

```

```
Out[20]:
```

	sex	length	diameter	height	whole_weight	shucked_weight	viscera_weight	shell_weight	age
0	2	0.455	0.365	0.095	0.5140	0.2245	0.1010	0.150	16.5
1	2	0.350	0.265	0.090	0.2255	0.0995	0.0485	0.070	8.5
2	0	0.530	0.420	0.135	0.6770	0.2565	0.1415	0.210	10.5
3	2	0.440	0.365	0.125	0.5160	0.2155	0.1140	0.155	11.5
4	1	0.330	0.255	0.080	0.2050	0.0895	0.0395	0.055	8.5

```
In [21]: x=df.iloc[:, :-1]
x.head()
```

```
Out[21]:
```

	sex	length	diameter	height	whole_weight	shucked_weight	viscera_weight	shell_weight
0	2	0.455	0.365	0.095	0.5140	0.2245	0.1010	0.150
1	2	0.350	0.265	0.090	0.2255	0.0995	0.0485	0.070
2	0	0.530	0.420	0.135	0.6770	0.2565	0.1415	0.210
3	2	0.440	0.365	0.125	0.5160	0.2155	0.1140	0.155

```
Out[21]:
```

	sex	length	diameter	height	whole_weight	shucked_weight	viscera_weight	shell_weight
0	2	0.455	0.365	0.095	0.5140	0.2245	0.1010	0.150
1	2	0.350	0.265	0.090	0.2255	0.0995	0.0485	0.070
2	0	0.530	0.420	0.135	0.6770	0.2565	0.1415	0.210
3	2	0.440	0.365	0.125	0.5160	0.2155	0.1140	0.155
4	1	0.330	0.255	0.080	0.2050	0.0895	0.0395	0.055

```
In [24]: y=df.iloc[:, -1]
y.head()
```

```
Out[24]:
```

0	16.5
1	8.5
2	10.5
3	11.5
4	8.5

Name: age, dtype: float64

```
In [22]: from sklearn.preprocessing import StandardScaler
scaler=StandardScaler()
x=scaler.fit_transform(x)
```

```
In [25]: x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.33)
```

```
In [26]: x_train.shape
```

```
Out[26]: (2797, 8)
```

```
In [27]: x_test.shape
```

```
Out[27]: (1378, 8)
```

```
In [28]: from sklearn.ensemble import RandomForestRegressor
reg=RandomForestRegressor()
```

In [29]: `reg.fit(x_train,y_train)`

Out[29]: `RandomForestRegressor()`

In [30]: `y_pred=reg.predict(x_test)`

In [31]: `from sklearn.metrics import mean_squared_error`
`import math`
`print(math.sqrt(mean_squared_error(y_test,y_pred)))`

1.8306598073341425